

Pastlens training loss: 5.9 $\rightarrow$ $\sim$ 2.3, most of the drop in the first $\sim$ 300 steps
1 epoch, 776k on-policy pairs, eff batch 64, lr 1e-4 (rsLoRA r64/ α 16 all-modules)

